

MIT EECS 6.815/6.865: Assignment 10:

Make your own assignment

Due Wednesday December 9 at 9pm

1 Summary

You'll choose your own assignment for this last one. This should be about the difficulty of an average pset.

You can choose a pset from a previous year that was not covered this year as long as you do some of the associated extra credit or extend it in some way (see <http://stellar.mit.edu/S/course/6/sp11/6.815/homework/index.html> and <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/index.html>).

Or you can create your own pset possibly from the list below. In addition to completing the assignment, you'll need to produce a short (2-4 page) writeup with the sections described in the deliverables section.

Make your own assignment! Turn in your writeup in PDF format to Stellar. Turn in your code to the submission system. See below for details. Choose one from a previous year or one from below or literally do your own thing. If you choose something not from a previous year and not on the list below, email us, so we can tell you if your idea is reasonable.

2 Deliverables

Please turn in a PDF file (you will lose points for any other format) describing your work and showing results. Give a short description (half a page to a page) of the algorithm that could be understood by someone who has taken the class but has not heard of that particular technique before. In addition, make sure your writeup has at least the following information in it.

- If you choose a pset from a previous year or a pset on the list below, indicate its name clearly.
- Describe what you implemented, what it does and what was difficult about implementing it.
- Background section with a summary of at least 3 papers related to what you are working on.
- At least 5 well-documented test cases testing intermediate parts of your algorithm. If possible add figures to the writeup about your test cases. This should be similar to the test cases we have provided you on previous assignments. Include your test cases in the `a10_main.cpp` file.

- Your algorithm run on at least two different inputs including at least one you created. The more the better. If possible, add figures to your writeup showing the inputs and the results.
- Include running times and other stats when appropriate. In general, anything that sheds light on the technique and its performance is good.
- The total document should probably be about 2-4 pages.
- Include your name at the beginning.

In addition to your writeup, please include your code. We provide you with a skeleton zip file containing some code from previous assignments and a makefile. Add your test cases to `a10_main.cpp` and your code and functions to `a10_pset.cpp` and `a10_pset.h`. Feel free to add additional source and header files. If you do, you'll need to add them to the `Makefile`. You'll need to put your function signature in the `.h` file. See previous assignments for details on how to do that. Submit the code to the submission system as with assignments 0 through 9 and make sure that your code runs on it. If you have any trouble or feel like the submission system is not sufficient for your needs (e.g. you need to use an external library or you want to use Halide), let us know.

6.865 students To get full credit for 6.865, you need to implement a little more than what is described below or provide additional analysis. In the rest of the document, when we say “to get full credit”, we mean to get full credit in 6.815, except in the “harder” section. In many cases, additional components are described and you can just pick from there. It's even easier for assignments from previous offerings: just do the grad version. In general, if you're unsure, email us.

3 Previous Years' Assignments

- Bayesian Matting: <http://stellar.mit.edu/S/course/6/sp11/6.815/homework/assignment5/>
- Seam Carving: <http://stellar.mit.edu/S/course/6/sp11/6.815/homework/assignment6/>
- Deconvolution and Poisson Editing <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment10/>
- Non-photorealistic Rendering <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment11/>
- Light Field (Lytro): <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment12/>

4 Easy

4.1 Texture synthesis

Given input texture example, generate a similar-looking but potentially bigger texture.

<http://graphics.cs.cmu.edu/people/efros/research/EfrosLeung.html>
http://en.wikipedia.org/wiki/Texture_synthesis
<http://www.ics.uci.edu/~fowlkes/class/cs116/hwk3/index.html>
<http://www.cs.ubc.ca/~woodham/cpsc425/assignments/hw4/hw4.html>
http://cs.nyu.edu/~fergus/teaching/comp_photo/assign3.pdf

For full credit, perform hole filling.

4.2 Flash no flash photography

Implement Petschnigg's version first, it is simpler because it doesn't seek to deal with shadows.

<http://dl.acm.org/citation.cfm?id=1015777> <http://people.csail.mit.edu/fredo/PUBLI/flash/index.htm>

Very similar to the tone mapping assignment.

Just implementing Petschnigg's approach won't give you full credit. For that, implement either a shadow fix or an alignment procedure.

Data is available on Elmar's page <http://maverick.inria.fr/Publications/2004/ED04/index.php> but we highly encourage you to capture your own. You can even borrow a camera that takes a flash and a no-flash image in succession.

4.3 Hybrid images

Generate images that look different from a close vs. a large distance. <http://cvcl.mit.edu/hybridimage.htm>

http://www.cs.illinois.edu/class/fa11/cs498dh/projects/hybrid/ComputationalPhotography_ProjectHybrid.html

To get full credit, show a plot of the frequency content of the two input images, the hybrid image, and its two components, and experiment with color.

Include at least two examples.

You might want to use your warping code to align the two images.

5 Normal, with instructions

5.1 Image resizing using seam carving

Make an image smaller by removing pixels that are less significant.

A little bit of Sobel gradient to create an energy function that says which pixels can be removed, then standard dynamic programming to find "seams" that connect easily-removable pixels.

<http://www.faculty.idc.ac.il/arik/site/seam-carve.asp>

<https://stellar.mit.edu/S/course/6/sp11/6.815/homework/assignment6/>

5.2 Bayesian matting

See <https://stellar.mit.edu/S/course/6/sp11/6.815/homework/assignment5/>

5.3 Inpainting with big database

Replace a masked area in an image by content found in a similar image from a big database of pictures.

<http://www.cs.brown.edu/courses/csci1950-g/asgn/proj4/>
<http://www.cs.brown.edu/courses/csci1950-g/asgn/proj4/resources/SceneCompletion.pdf>

5.4 Tour into the Picture

Create 3D animations from a single image and a few clicks!

http://graphics.cs.cmu.edu/courses/15-463/2007_fall/Papers/TIP.pdf
Just start with your homography code. It's OK if you have to manually indicate where the four corners should go. Then add the notion of vertical billboard.
http://graphics.cs.cmu.edu/courses/15-463/2010_fall/hw/proj4g/

6 Normal

6.1 Dehazing

Remove haze in photography. Locally compute the minimum (scipy has a min filter) and use it to estimate what to subtract from the image.

http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=5206515&tag=1

Ignore the soft matting from that paper. Replace it by a cross-bilateral filter, which is easier to implement.

6.2 Deconvolution (easy to implement, requires a little bit of math to understand)

For this one, you may want to refer to <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment10/>.

Given a blurry image, invert the blur process to yield a sharp image. If the blur process is described by the convolution operator A and your input blurry image is y , you want to solve for $Ax = b$, which is very similar to the Poisson equation we solved. You only need to replace the Laplacian operator by the blur kernel.

Make the process more stable by adding a gradient-based regularization. To avoid amplifying the noise, minimize:

$$\min ||Ax - y||^2 + \lambda ||\nabla x||^2$$

where λ is a parameter.

Extra-credit (easy to implement, hard to understand): Use reweighted least square to simulate an L1 regularization. That is, rather than minimizing the gradients with uniform weights, reweight the constraint of each gradient by the magnitude of the gradient.

Careful: you need to reweight the gradient, not the Laplacian. You need to decompose the Laplacian as the divergence of the gradient. As discussed in class, to compute the divergence of the gradient, you should use kernels $[-1, 1]$ but you need to use forward difference for one and backward differences for the other one, so that the overall kernel gets centered.

Or implement the Richardson-Lucy version http://en.wikipedia.org/wiki/Richardson%E2%80%93Lucy_deconvolution

6.3 Stainglass by numbers

For this assignment, you may want to refer to <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment11/>.

Similar to the painterly rendering, except that we store some notion of depth for each pixel (z buffer) and splat 3D cones centered at y, x .

That is, store an array of z values. For each brush stroke, the idea is to render a 3D cone center at the stroke and extending away from the canvas. That is, the z value for a pixel is equal to its distance from the center of the stroke. Only the closest cone is visible at a point, so you need to test for each pixel if the new cone is closer than the stored value. If yes, update both your z array and the color array.

Vary the narrowness of the regions by applying a different scale factor to compute depth for each cone.

<http://dl.acm.org/citation.cfm?id=97902>

Forget about the relaxation part, unless you're very motivated.

6.4 Denoising by wavelet coring

<http://www.cns.nyu.edu/pub/lcv/simoncelli96c.pdf>

6.5 Video texture

Create videos that loop perfectly, and even graphs of transition for non-repetitive playing.

<http://www.cc.gatech.edu/cpl/projects/videotexture/SIGGRAPH2000/index.htm>

6.6 Color 2 gray

Turn a color image into a black and white one while preserving edges as well as possible. Start from your Poisson code and the max of the gradient across the three channels. Then be smarter. <http://www.cs.northwestern.edu/~ago820/color2gray/>

6.7 NL means denoising (easy but slow)

<http://bengal.missouri.edu/~kes25c/nl1.pdf> http://www.ipol.im/pub/algo/bcm_non_local_means_denoising/

6.8 Morphable face models and caricatures

<http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.49.9275>
http://web.mit.edu/emeyers/www/face_databases.html
<http://vasc.ric.cmu.edu/idb/html/face/>

Start with your morphing code and create the same segments for a lot of different faces. I can give you what I have from the pset for all the students, but I wasn't fully happy with the results. It might work better with one of the standard datasets I linked, because people are all in exactly the same pose and their hair usually doesn't get as much in the way.

Things I would like to see:

- average face
- average male, average female
- caricature of a face.
- make a face more male or more female

6.9 Pyramid image alignment

Implement a coarse-to-fine version of image alignment.

Extend it to the median pyramid by Greg Ward http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.59.80%93Kanade_method

Or go full Lucas-Kanade, e.g. http://en.wikipedia.org/wiki/Lucas%E2%80%93Kanade_method

6.10 Time lapse manipulation

<http://dl.acm.org/citation.cfm?id=1276505> or http://www.csbio.unc.edu/mcmillan/pubs/sig07_Bennett.pdf

Median or min across time, dynamic programming.

For full credit, implement at least the dynamic programming version with two different metrics.

6.11 Patch match

Called content-aware fill in Photoshop.

http://gfx.cs.princeton.edu/pubs/Barnes_2009_PAR/

6.12 Detecting copy-pasting

http://www.cs.dartmouth.edu/farid/Hany_Farid/Papers/Entries/2011/3/20_Exposing_Digital_Forgeries_by_Detecting_Duplicated_Image_Regions.html

To get full credit, extend to detecting Poisson image cloning. Easy but probably slow. Try to accelerate using convolution/correlation.

6.13 Photographic style transfer

<http://people.csail.mit.edu/soonmin/photolook/>

Focus on histogram matching of the bilateral filter components, and in particular the notion of texture, which is not very different from the sharpness in pset 11. Don't worry too much about the post-processing and the gradient preservation unless you have time.

You'll get full credit if you get the transfer of global contrast and texture. Post-processing effects and gradient preservation are extra credit (unless your in 6.865).

6.14 Salomon-style art

Jason Salomon does amazing algorithmic art, usually based on the combination of many photos. <http://salomon.com/work/>. See also <http://blog.xkcd.com/2010/05/03/color-survey-results/>

Use the flickr API <http://www.flickr.com/services/api/> to reproduce images such as his color wheel <http://salomon.com/work/color-wheel/image/409/> by querying flickr for images with color names such as "red". Start with a flat rectangular version. See http://en.wikipedia.org/wiki/Color_term for more inspiration and create a multilingual comparison.

Aggregate many photos of a given landmark ("Statue of Liberty") or type of image ("landscape") or ("portrait") in the spirit of <http://salomon.com/work/Homes/grid/2/>, <http://salomon.com/work/Portrait/grid/1/>. Maybe cluster the results somehow to create multiple composites.

You can always do old-style image mosaics, but at least try to match the edge structure of the individual super pixel images to that of the target photo. http://en.wikipedia.org/wiki/Photographic_mosaic

To get full credit, create at least two of these (e.g. the color wheel and the landmark), or one with extra bells and whistle (e.g. landmark+clustering, multilingual color wheel).

6.15 Anisotropic diffusion

http://en.wikipedia.org/wiki/Anisotropic_diffusion

Alternative to the bilateral filter, but based on PDEs.

6.16 Laplacian pyramids

The generalization of the 2-scale blending that we did, which can also be used for the apple/orange trick or the focal stack fusion.

http://persci.mit.edu/pub_pdfs/RCA85.pdf http://www.cs.princeton.edu/courses/archive/spr04/cos429/papers/burt_adelson.pdf

6.17 Photobios

Implement photobios from <http://grail.cs.washington.edu/photobios/paper.pdf>, <http://grail.cs.washington.edu/photobios/video.mp4>. You can try getting data from this website: <http://daily.jasonfletcher.info/>. Or find your own collection of lots of faces.

7 Normal, but requires hardware

7.1 Separation of direct and indirect lighting effects

<http://www1.cs.columbia.edu/CAVE/projects/separation/>

You can borrow a projector.

7.2 Dual photography

http://graphics.stanford.edu/papers/dual_photography/

You can borrow a projector.

7.3 Relighting with multiple photographs

Take images with a static camera (on tripod) but with light coming from different directions. Then use these images to create new images using weighted combinations. See e.g. <http://gl.ict.usc.edu/Research/LS3/> for inspiration, but don't try to reproduce all their crazy stuff.

8 Harder

8.1 View morphing

Combine homographies and morphing! Add a homography to make your morphing respect 3D structure better. In particular, given two views of the same 3D object, this method guarantees that the morphing sequence is equal to a 3D rotation around the object.

<http://www.cs.washington.edu/homes/seitz/papers/sigg96.pdf>

The paper is not completely easy to read but it's cool.

8.2 Lens correction

Calibration and correction of radial distortion and vignetting. See the slides.

Vignetting is more tricky than it seems. replace the polynomial by a table and a piecewise-linear function.

8.3 Inpainting (not so hard to implement but not easy to understand)

For this assignment, you may want to refer to <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment10/>.

Given an image and a masked region, reconstruct plausible values inside the mask by interpolation. <http://en.wikipedia.org/wiki/Inpainting> <http://www.tecn.upf.es/~mbertalmio/restoration0.html> http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=5593835

I suggest you combine the Poisson solver and the structure tensor. First compute the structure tensor, ignoring pixels in the masked region. (The local weighted average in the second Gaussian blur should ignore pixels in the region. The easiest way to do it is to set them to zero, and keep track of the sum of the weights in the sum by blurring the mask itself.). Then run Poisson with a flat source to interpolate the structure tensor inside the region. Once you have an interpolated structure tensor, use it to interpolate color values, for example using anisotropic diffusion http://en.wikipedia.org/wiki/Anisotropic_diffusion, where the diffusion is guided by the 2x2 structure tensor matrix at each pixel.

8.4 Bundle adjustment

Refine your panorama with a global optimization, including radial distortion optimization. This one might be really hard in C++ because of the optimization. You can try using <http://ab-initio.mit.edu/wiki/index.php/NLopt> for optimization.

<http://sse.tongji.edu.cn/linzhang/computervision/projects/image%20alignment%20and%20stitching%20a%20tutorial.pdf>

<http://www.cs.jhu.edu/~misha/ReadingSeminar/Papers/Triggs00.pdf>

Given your inlier pairs for the various images you have, optimize free parameters to minimize the reproduction error. Start by writing this error function computation. Make it “robust” by clamping it to a max value (if the reproduction is more than XXX pixels away, only pay the penalty for XXX pixels).

The free parameters are typically the focal length and three 3D rotation angles for each photo, 3D coordinates for each point (of course up to scale, since we can't know the distance to the camera), and some radial distortion parameter (which you should forget about at the beginning). Good initialization is critical.

Start from your homographies, guess focal length (e.g. 30mm for a 24x36mm sensor), and look at the maths of projection from the cylindrical panorama assignment.

Start with a brute force approach. This will be enough to get full credit. Demonstrate that your method can, e.g. converge to the correct focal length (use a pano where you know the focal length, but initialize with a wrong one). Your first read should probably be section 5.1 of Szeliski's survey above.

http://web.me.com/dellaert/07F-Vision/Schedule_files/08-Stitching.ppt.pdf

<http://research.google.com/pubs/pub37112.html>

8.5 Colorization using least square optimization

For this assignment, you may want to refer to <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment10/>.

Given a greyscale image and a sparse set of color indications given by the user, propagate these colors to the full image. The interpolation takes into account the content of the greyscale image and tends to have color changes only where the intensity changes.

See <http://www.cs.huji.ac.il/~yweiss/Colorization/>

You can do a full linear algebra version by forming the sparse matrix and use your Poisson code but with BF

You can alternatively modify your Poisson image editing code and replace the Laplacian kernel by an input-dependent kernel.

8.6 Perceptual metric for photo retouching

http://www.cs.dartmouth.edu/farid/Hany_Farid/Papers/Entries/2011/6/6_A_Perceptual_Metric_for_Photo_Retouching.html

Use your morphing code for computing the warp field.

The tricky part is to get data.

8.7 Hockney collage from multiple images

<http://webee.technion.ac.il/~lihi/Demos/AutoJoiners.html>

Modify your automatic panorama matching and perform automatic layout using least square optimization, trying to use the average translation vector between pairs of images.

Speed could be an issue.

8.8 Hockney collage from a single image

Create a collage in the style of David Hockney http://www.hockneypictures.com/works_photos.php

See an example of software at <http://bighugelabs.com/hockney.php>

I'd start with the NPR algorithm to scatter a bunch of window locations across the image according to the importance map.

8.9 Local Laplacian

What replaced the bilateral filter in Camera RAW/Lightroom.

<http://people.csail.mit.edu/sparis/publi/2011/siggraph/>

8.10 Image deformation using moving least squares

Related to warping. Specify a sparse set of point displacement and interpolate intelligently by solving a least-square problem at each point. In particular, it allows the interpolation to have some underlying notion of class of transformations, such as angle preservation.

Probably slow. Use numpy to solve each least square problem.

<http://faculty.cs.tamu.edu/schaefer/research/mls.pdf>

An extension uses biharmonic energies

<http://igl.ethz.ch/projects/bbw/>

8.11 Non-Photorealistic rendering using extended differences of Gaussians

For this one, you may want to refer to <https://stellar.mit.edu/S/course/6/fa13/6.815/homework/assignment11/>.

<http://dl.acm.org/citation.cfm?id=2024700>

Use the structure tensor and the eigenvector business from pset 12 to get the flow alignment, or use the Sobel operator if you prefer.

Implement both adaptive thresholding and flow alignment and you'll get full credit.

Rather than being circular, Gaussians now have elliptical shapes, and the long axis of the Gaussian follows the directions from the tensor field. To get an anisotropic Gaussian, rather than using $\exp(-(x^2 + y^2)/2\sigma^2)$ you replace the simple square by a general quadratic: $\exp(-V^T S V)$ where V is your vector x, y and S is a symmetric positive definite matrix. For example, the diagonal matrix $\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$ gives you a Gaussian elongated along the second coordinate.

One way to achieve general anisotropic Gaussians is to generate a canonical one like the one just above and then use the image rotation function we used in pset 12 to get other orientations.

See also <http://www.science.uva.nl/research/publications/2003/GeusebroekTIP2003/>

8.12 Multiflash camera

<http://web.media.mit.edu/~raskar/NprCamera/>

8.13 Guided image filtering

<http://research.microsoft.com/en-us/um/people/kahe/eccv10/eccv10supp.pdf>

8.14 graph cut / grab cut

Foreground/background extraction

<http://www.csd.uwo.ca/~yuri/Abstracts/iccv01-abs.html> <http://research.microsoft.com/apps/pubs/default.aspx?id=67890>

Implement graph cut segmentation with a combination of data and edge term and you'll get full credit. Grab cut is extra credit. Don't worry about the mixture of Gaussian part and keep the same histogram approach.

8.15 Interactive digital photomontage

<http://grail.cs.washington.edu/projects/photomontage/>

8.16 Reducing veiling glare for higher-dynamic-range imaging

http://graphics.stanford.edu/papers/glare_removal/

8.17 Artistic screening

http://www.iro.umontreal.ca/~ostrom/publications/pdf/SIGGRAPH95_ArtisticScreening.pdf reproduce shades of grey with micro patterns of your choice.

really hardcore: color version http://www.iro.umontreal.ca/~ostrom/publications/pdf/SIGGRAPH99_MultiColorDithering_600dpi.pdf

8.18 Laplacian matting

Separate foreground and background. The derivation is a little scary but the implementation can be simplish. <http://www.wisdom.weizmann.ac.il/~levina/papers/Matting-Levin-Lischinski-Weiss-CVPR06.pdf>

8.19 Inpainting with parametric texture synthesis

<http://people.csail.mit.edu/torralba/courses/6.869/lectures/lecture5/heegerbergen.pdf>

http://graphics.stanford.edu/papers/texture_replace/

8.20 More inpainting

http://www.itc.ku.edu/~potetz/EECS_841_Fall08/Readings/Lecture_39_CriminisiPerezToyama_Inpainting_04.pdf